

```
mkdir grade
```

```
nano StudentGrade.java
```

```
nano students.txt
```

```
javac -classpath `hadoop classpath` -d output StudentGrade.java
```

```
jar -cvf studentgrade.jar -C output .
```

```
hdfs dfs -mkdir /user/hadoop/output_studentgrade
```

```
hdfs dfs -put students.txt /user/hadoop/input
```

```
hadoop jar studentgrade.jar StudentGrade /user/hadoop/input /user/hadoop/output_studentgrade
```

```
hdfs dfs -cat /user/hadoop/output_studentgrade/part-r-00000
```

Exp 2

Java file :

```
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.fs.Path;
```

```
import org.apache.hadoop.io.IntWritable;
```

```
import org.apache.hadoop.io.Text;
```

```
import org.apache.hadoop.mapreduce.Job;
```

```
import org.apache.hadoop.mapreduce.Mapper;
```

```
import org.apache.hadoop.mapreduce.Reducer;
```

```
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```
import java.io.IOException;
```

```
import java.util.StringTokenizer;
```

```
public class WordCount {
```

```
    public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable> {
```

```
        private final static IntWritable one = new IntWritable(1);
```

```
        private Text word = new Text();
```

```
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
```

```
            StringTokenizer itr = new StringTokenizer(value.toString());
```

```
            while (itr.hasMoreTokens()) {
```

```
                word.set(itr.nextToken());
```

```

        context.write(word, one);
    }
}

}

public static class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {

        int sum = 0;

        for (IntWritable val : values) {

            sum += val.get();

        }

        result.set(sum);

        context.write(key, result);

    }

}

public static void main(String[] args) throws Exception {

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "word count");

    job.setJarByClass(WordCount.class);

    job.setMapperClass(TokenizerMapper.class);

    job.setCombinerClass(IntSumReducer.class);

    job.setReducerClass(IntSumReducer.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}

```

Exp 3

Java :

```
import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;


import java.io.IOException;


public class StudentGrade {


    public static class GradeMapper extends Mapper<Object, Text, Text, Text> {

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] parts = value.toString().split(",");

            if (parts.length == 2) {

                String studentName = parts[0].trim();

                double score = Double.parseDouble(parts[1].trim());

                String grade;

                if (score >= 90) {

                    grade = "A";

                } else if (score >= 80) {

                    grade = "B";

                } else if (score >= 70) {

                    grade = "C";

                } else if (score >= 60) {

                    grade = "D";

                } else {

                    grade = "F";

                }

                context.write(new Text(studentName), new Text(grade));

            }

        }

    }

}
```

```
}
```

```
public static class GradeReducer extends Reducer<Text, Text, Text, Text> {  
  
    public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {  
  
        for (Text val : values) {  
  
            context.write(key, val);  
  
            break; // Only one value per student expected  
  
        }  
  
    }  
  
}
```

```
public static void main(String[] args) throws Exception {  
  
    if (args.length != 2) {  
  
        System.err.println("Usage: StudentGrade <input path> <output path>");  
  
        System.exit(-1);  
  
    }  
  
}
```

```
Configuration conf = new Configuration();  
  
Job job = Job.getInstance(conf, "Student Grade Calculation");
```

```
job.setJarByClass(StudentGrade.class);  
  
job.setMapperClass(GradeMapper.class);  
  
job.setReducerClass(GradeReducer.class);  
  
job.setOutputKeyClass(Text.class);  
  
job.setOutputValueClass(Text.class);
```

```
FileInputFormat.addInputPath(job, new Path(args[0]));  
  
FileOutputFormat.setOutputPath(job, new Path(args[1]));
```

```
System.exit(job.waitForCompletion(true) ? 0 : 1);  
  
}  
  
}
```

Exp 4

Java

```
import org.apache.hadoop.conf.Configuration;  
  
import org.apache.hadoop.fs.Path;  
  
import org.apache.hadoop.io.DoubleWritable;
```

```

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;


import java.io.IOException;


// Main class

public class MaxTemperature {


    // Mapper Class

    public static class TemperatureMapper extends Mapper<Object, Text, IntWritable, DoubleWritable> {

        private IntWritable year = new IntWritable();

        private DoubleWritable temperature = new DoubleWritable();


        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] fields = value.toString().split(","); // Assuming CSV format

            if (fields.length < 3) return; // Skip malformed lines


            try {

                year.set(Integer.parseInt(fields[0].trim())); // Extract Year

                temperature.set(Double.parseDouble(fields[2].trim())); // Extract Temperature

                context.write(year, temperature);

            } catch (NumberFormatException e) {

                // Skip invalid data

            }

        }

    }


    // Reducer Class

    public static class TemperatureReducer extends Reducer<IntWritable, DoubleWritable, IntWritable, DoubleWritable> {

        public void reduce(IntWritable key, Iterable<DoubleWritable> values, Context context)

            throws IOException, InterruptedException {

            double maxTemp = Double.MIN_VALUE;

            for (DoubleWritable val : values) {

```

```

        maxTemp = Math.max(maxTemp, val.get());
    }

    context.write(key, new DoubleWritable(maxTemp));
}
}

// Driver Code

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Max Temperature Finder");

    job.setJarByClass(MaxTemperature.class);
    job.setMapperClass(TemperatureMapper.class);
    job.setReducerClass(TemperatureReducer.class);
    job.setMapOutputKeyClass(IntWritable.class);
    job.setMapOutputValueClass(DoubleWritable.class);
    job.setOutputKeyClass(IntWritable.class);
    job.setOutputValueClass(DoubleWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0])); // Input path
    FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output path

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Exp 5

Java:

```

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

```

```

import java.util.ArrayList;

import java.util.List;

public class MatrixMultiplication {

    // Mapper Class

    public static class MatrixMapper extends Mapper<Object, Text, Text, Text> {

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] parts = value.toString().split("\\s+");

            if (parts.length < 4) return;

            String matrixName = parts[0]; // 'A' or 'B'

            int i = Integer.parseInt(parts[1]);

            int j = Integer.parseInt(parts[2]);

            double val = Double.parseDouble(parts[3]);

            Configuration conf = context.getConfiguration();

            int p = Integer.parseInt(conf.get("p")); // Columns of B

            if (matrixName.equals("A")) {

                for (int k = 0; k < p; k++) {

                    context.write(new Text(i + "," + k), new Text("A," + j + "," + val));

                }

            } else if (matrixName.equals("B")) {

                int m = Integer.parseInt(conf.get("m")); // Rows of A

                for (int k = 0; k < m; k++) {

                    context.write(new Text(k + "," + j), new Text("B," + i + "," + val));

                }

            }

        }

    }

    // Reducer Class

    public static class MatrixReducer extends Reducer<Text, Text, Text, Text> {

        public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {

            List<String> aValues = new ArrayList<>();

            List<String> bValues = new ArrayList<>();

        }

    }

```

```

    for (Text val : values) {

        String valueStr = val.toString();

        if (valueStr.startsWith("A")) {

            aValues.add(valueStr);

        } else {

            bValues.add(valueStr);

        }

    }

}

double sum = 0;

for (String a : aValues) {

    String[] aParts = a.split(",");

    int aCol = Integer.parseInt(aParts[1]);

    double aVal = Double.parseDouble(aParts[2]);

    for (String b : bValues) {

        String[] bParts = b.split(",");

        int bRow = Integer.parseInt(bParts[1]);

        double bVal = Double.parseDouble(bParts[2]);

        if (aCol == bRow) {

            sum += aVal * bVal;

        }

    }

}

context.write(key, new Text(Double.toString(sum)));

}

}

// Main Method

public static void main(String[] args) throws Exception {

    Configuration conf = new Configuration();

    conf.set("m", "2"); // Rows of A

    conf.set("n", "3"); // Columns of A / Rows of B

    conf.set("p", "2"); // Columns of B

    Job job = Job.getInstance(conf, "Matrix Multiplication");

```



```

        job.setJarByClass(MatrixMultiplication.class);
        job.setMapperClass(MatrixMapper.class);
        job.setReducerClass(MatrixReducer.class);

        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(Text.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(Text.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

Exp 6

Java:

```

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

public class MaxElectricityConsumption {

    // Mapper Class
    public static class MaxConsumptionMapper extends Mapper<Object, Text, Text, IntWritable> {

        private Text year = new Text();
        private IntWritable consumption = new IntWritable();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

```

```

String[] fields = value.toString().split(",");

if (fields.length == 3) {

    year.set(fields[0].trim()); // Extract Year

    consumption.set(Integer.parseInt(fields[2].trim())); // Extract Consumption

    context.write(year, consumption);

}

}

}

// Reducer Class

public static class MaxConsumptionReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

        int maxConsumption = Integer.MIN_VALUE;

        for (IntWritable val : values) {

            maxConsumption = Math.max(maxConsumption, val.get());

        }

        context.write(key, new IntWritable(maxConsumption));

    }

}

// Driver Code

public static void main(String[] args) throws Exception {

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Max Electricity Consumption");

    job.setJarByClass(MaxElectricityConsumption.class);

    job.setMapperClass(MaxConsumptionMapper.class);

    job.setReducerClass(MaxConsumptionReducer.class);

    job.setMapOutputKeyClass(Text.class);

    job.setMapOutputValueClass(IntWritable.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));

    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);

```

```
}  
}
```

Exp 7

Java:

```
import java.io.IOException;  
  
import java.util.StringTokenizer;  
  
  
import org.apache.hadoop.conf.Configuration;  
  
import org.apache.hadoop.fs.Path;  
  
import org.apache.hadoop.io.IntWritable;  
  
import org.apache.hadoop.io.Text;  
  
import org.apache.hadoop.mapreduce.Job;  
  
import org.apache.hadoop.mapreduce.Mapper;  
  
import org.apache.hadoop.mapreduce.Reducer;  
  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
  
  
public class MaxElectricalConsumption {  
  
    // Mapper class to extract month and consumption data  
    public static class ConsumptionMapper extends Mapper<Object, Text, Text, IntWritable> {  
  
        private Text month = new Text();  
  
        private IntWritable consumption = new IntWritable();  
  
  
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {  
  
            StringTokenizer itr = new StringTokenizer(value.toString());  
  
            // Expected format: "Month Consumption"  
  
            if (itr.hasMoreTokens()) {  
  
                String monthString = itr.nextToken(); // Month (e.g., "January")  
  
                if (itr.hasMoreTokens()) {  
  
                    String consumptionString = itr.nextToken(); // Consumption value  
  
                    try {  
  
                        month.set(monthString);  
  
                        consumption.set(Integer.parseInt(consumptionString));  
  
                        context.write(month, consumption);  
  
                    } catch (NumberFormatException e) {  
  
                        System.err.println("Skipping invalid record: " + value.toString());  
  
                    }  
  
                }  
  
            }  
  
        }  
  
    }  
  
}
```

```

    }
}
}
}
}

```

// Reducer class to calculate the maximum consumption across all months

```

public static class MaxConsumptionReducer extends Reducer<Text, IntWritable, Text, IntWritable> {

    private Text resultMonth = new Text();

    private IntWritable resultConsumption = new IntWritable();

    private String maxMonth = "";

    private int maxConsumption = Integer.MIN_VALUE;

    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException, InterruptedException {

        for (IntWritable val : values) {

            if (val.get() > maxConsumption) {

                maxConsumption = val.get();

                maxMonth = key.toString();

            }

        }

    }

}

```

@Override

```

protected void cleanup(Context context) throws IOException, InterruptedException {

    resultMonth.set(maxMonth);

    resultConsumption.set(maxConsumption);

    context.write(resultMonth, resultConsumption);

}

}

```

// Driver code

```

public static void main(String[] args) throws Exception {

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Maximum Electrical Consumption");

    job.setJarByClass(MaxElectricalConsumption.class);

    job.setMapperClass(ConsumptionMapper.class);

    job.setReducerClass(MaxConsumptionReducer.class);

```

```

        job.setOutputKeyClass(Text.class); // Month
        job.setOutputValueClass(IntWritable.class); // Consumption

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

Exp 8

```

import java.io.IOException;
import java.util.ArrayList;
import java.util.List;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MovieTagsAnalysis {

    // Mapper Class

    public static class TagsMapper extends Mapper<Object, Text, Text, Text> {

        private Text movieID = new Text();
        private Text tag = new Text();

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] fields = value.toString().split(",");

            if (fields.length >= 3) {

                movieID.set(fields[1]); // Movie ID is in the second column
            }
        }
    }
}

```

```

        tag.set(fields[2]); // Tag is in the third column

        context.write(movieID, tag);
    }
}
}

// Reducer Class

public static class TagsReducer extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {

        List<String> tagsList = new ArrayList<>();

        for (Text val : values) {

            tagsList.add(val.toString());

        }

        context.write(key, new Text(String.join(" ", tagsList)));

    }
}

// Driver Code

public static void main(String[] args) throws Exception {

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Movie Tags Analysis");

    job.setJarByClass(MovieTagsAnalysis.class);

    job.setMapperClass(TagsMapper.class);

    job.setReducerClass(TagsReducer.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0])); // Input path (e.g., tags.csv)

    FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output path

    System.exit(job.waitForCompletion(true) ? 0 : 1);

}
}

```

Exp 9

```
import java.io.IOException;
```

```

import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;


public class UberTripsAnalysis {


    // Mapper Class

    public static class UberMapper extends Mapper<Object, Text, Text, Text> {

        private Text baseId = new Text();

        private Text dateAndTrips = new Text();


        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] tokens = value.toString().split("\\s+"); // Split by whitespace

            if (tokens.length == 4) { // Expected: base_id date vehicles trips

                String base = tokens[0]; // Base ID

                String date = tokens[1]; // Date

                String trips = tokens[3]; // Trips (assumed at 4th column)


                baseId.set(base);

                dateAndTrips.set(date + "," + trips);

                context.write(baseId, dateAndTrips);

            }

        }

    }


    // Reducer Class

    public static class UberReducer extends Reducer<Text, Text, Text, Text> {

        private Text maxTripDate = new Text();


        public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {

            int maxTripsValue = Integer.MIN_VALUE;

            String maxDay = "";

```

```

        for (Text val : values) {

            String[] parts = val.toString().split(",");

            if (parts.length == 2) {

                String date = parts[0];

                int trips = Integer.parseInt(parts[1]);

                if (trips > maxTripsValue) {

                    maxTripsValue = trips;

                    maxDay = date;

                }

            }

        }

        maxTripDate.set(maxDay + " " + maxTripsValue);

        context.write(key, maxTripDate);

    }

}

// Driver Code

public static void main(String[] args) throws Exception {

    if (args.length < 2) {

        System.err.println("Usage: UberTripsAnalysis <input file> <output dir>");

        System.exit(1);

    }

    Configuration conf = new Configuration();

    Job job = Job.getInstance(conf, "Uber Trips Analysis");

    job.setJarByClass(UberTripsAnalysis.class);

    job.setMapperClass(UberMapper.class);

    job.setReducerClass(UberReducer.class);

    job.setMapOutputKeyClass(Text.class);

    job.setMapOutputValueClass(Text.class);

    job.setOutputKeyClass(Text.class);

    job.setOutputValueClass(Text.class);

```



```

FileInputFormat.addInputPath(job, new Path(args[0]));

FileOutputFormat.setOutputPath(job, new Path(args[1]));


System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}

```

Exp 10

```

import java.io.IOException;


import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;


public class TitanicAnalysis {


    // Mapper Class

    public static class TitanicMapper extends Mapper<Object, Text, Text, Text> {

        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] columns = value.toString().split(",");


            // Skip header

            if (columns[0].equals("PassengerId")) {

                return;

            }


            try {

                int survived = Integer.parseInt(columns[1]); // 0 or 1

                int pclass = Integer.parseInt(columns[2]); // 1, 2, or 3

                String sex = columns[5]; // male/female

                String ageStr = columns[6]; // Age (can be empty)


                // Emit age of deceased passengers grouped by gender

```

```

        if (survived == 0 && !ageStr.isEmpty()) {
            context.write(new Text("Age_" + sex), new Text(ageStr + ",1"));
        }

        // Emit survivors grouped by class
        if (survived == 1) {
            context.write(new Text("Class_" + pclass), new Text("1,1"));
        }

    } catch (Exception e) {
        System.err.println("Skipping invalid record: " + value.toString());
    }
}

// Reducer Class
public static class TitanicReducer extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {

        double totalAge = 0;

        int count = 0;

        for (Text val : values) {
            String[] parts = val.toString().split(",");

            if (parts.length == 2) {
                double value = Double.parseDouble(parts[0]);

                int occurrences = Integer.parseInt(parts[1]);

                totalAge += value;

                count += occurrences;
            }
        }

        // Output average age for deceased passengers
        if (key.toString().startsWith("Age_")) {
            double avgAge = (count == 0) ? 0 : totalAge / count;

            context.write(key, new Text(String.format("Average Age: %.2f", avgAge)));
        }

        // Output survivor count per class

```

```

        else if (key.toString().startsWith("Class_")) {
            context.write(key, new Text("Total Survivors: " + count));
        }
    }
}

// Main Method

public static void main(String[] args) throws Exception {
    if (args.length < 2) {
        System.err.println("Usage: TitanicAnalysis <input file> <output dir>");
        System.exit(1);
    }

    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Titanic Data Analysis");

    job.setJarByClass(TitanicAnalysis.class);
    job.setMapperClass(TitanicMapper.class);
    job.setReducerClass(TitanicReducer.class);

    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(Text.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Exp 11

```

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;

```

```

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;


public class MaxTemperature {


    // Mapper Class

    public static class MaxTempMapper extends Mapper<Object, Text, Text, Text> {


        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {

            String[] columns = value.toString().split("\\s+"); // Split by whitespace

            if (columns.length != 3) return; // Skip invalid lines


            String year = columns[0]; // Extract year

            String month = columns[1]; // Extract month

            String temp = columns[2]; // Extract temperature


            // Emit (year, month_temperature)

            context.write(new Text(year), new Text(month + "_" + temp));

        }

    }


    // Reducer Class

    public static class MaxTempReducer extends Reducer<Text, Text, Text, IntWritable> {


        public void reduce(Text key, Iterable<Text> values, Context context)

            throws IOException, InterruptedException {

            int maxTemp = Integer.MIN_VALUE;

            String maxMonth = "";


            // Find the month with the highest temperature for each year

            for (Text val : values) {

                String[] parts = val.toString().split("_");

                String month = parts[0];

                int temperature = Integer.parseInt(parts[1]);

```

```

        if (temperature > maxTemp) {
            maxTemp = temperature;
            maxMonth = month;
        }
    }

    // Emit (Year, Month MaxTemperature)
    context.write(new Text(key.toString() + " " + maxMonth), new IntWritable(maxTemp));
}

}

// Main Method to configure and run the MapReduce job
public static void main(String[] args) throws Exception {
    if (args.length < 2) {
        System.err.println("Usage: MaxTemperature <input file> <output dir>");
        System.exit(1);
    }

    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Max Temperature Finder");

    job.setJarByClass(MaxTemperature.class);
    job.setMapperClass(MaxTempMapper.class);
    job.setReducerClass(MaxTempReducer.class);

    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(Text.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0])); // Input file path
    FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output directory

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

Exp 12

```
import org.apache.hadoop.conf.Configuration;

import org.apache.hadoop.fs.Path;

import org.apache.hadoop.io.IntWritable;

import org.apache.hadoop.io.LongWritable;

import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.Mapper;

import org.apache.hadoop.mapreduce.Reducer;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;

import org.apache.hadoop.io.Writable;

import java.io.DataInput;

import java.io.DataOutput;
```

```
public class Aggregate {
```

```
    // Writable class to aggregate Titanic data
```

```
    public static class TitanicData implements Writable {
```

```
        int totalPassengers;
```

```
        int survived;
```

```
        double totalAge;
```

```
        public TitanicData() {
```

```
            this.totalPassengers = 0;
```

```
            this.survived = 0;
```

```
            this.totalAge = 0.0;
```

```
        }
```

```
        // Add passenger data to the aggregate
```

```
        public void addPassenger(int survived, double age) {
```

```
            totalPassengers++;
```

```
            if (survived == 1) {
```

```
                this.survived++;
```

```
            }
```

```
            if (age > 0) {
```

```

        this.totalAge += age;
    }
}

// Calculate the survival rate
public double getSurvivalRate() {
    if (totalPassengers == 0) return 0.0;
    return (double) survived / totalPassengers * 100;
}

// Calculate the average age of passengers
public double getAverageAge() {
    if (totalPassengers == 0) return 0.0;
    return totalAge / totalPassengers;
}

// Get the total number of passengers
public int getTotalPassengers() {
    return totalPassengers;
}

@Override
public void write(DataOutput out) throws IOException {
    out.writeInt(totalPassengers);
    out.writeInt(survived);
    out.writeDouble(totalAge);
}

@Override
public void readFields(DataInput in) throws IOException {
    totalPassengers = in.readInt();
    survived = in.readInt();
    totalAge = in.readDouble();
}

}

// Mapper class to process Titanic dataset and map data by passenger class
public static class TitanicMapper extends Mapper<LongWritable, Text, Text, TitanicData> {

```

```

private Text classKey = new Text();

public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {

    String line = value.toString();

    String[] tokens = line.split(",");

    try {

        String pclass = tokens[2];

        int survived = Integer.parseInt(tokens[1]);

        double age = tokens[5].equals("") ? 0 : Double.parseDouble(tokens[5]);

        classKey.set(pclass);

        TitanicData data = new TitanicData();

        data.addPassenger(survived, age);

        context.write(classKey, data);

    } catch (NumberFormatException | ArrayIndexOutOfBoundsException e) {

        // Handle invalid records gracefully

    }

}

// Reducer class to aggregate data by passenger class
public static class TitanicReducer extends Reducer<Text, TitanicData, Text, Text> {

    private Text result = new Text();

    public void reduce(Text key, Iterable<TitanicData> values, Context context) throws IOException, InterruptedException {

        TitanicData aggregatedData = new TitanicData();

        for (TitanicData data : values) {

            aggregatedData.totalPassengers += data.getTotalPassengers();

            aggregatedData.survived += data.survived;

            aggregatedData.totalAge += data.totalAge;

        }

        double survivalRate = aggregatedData.getSurvivalRate();

        double avgAge = aggregatedData.getAverageAge();

        int totalPassengers = aggregatedData.getTotalPassengers();

        String resultString = String.format("Total Passengers: %d, Survival Rate: %.2f%%, Average Age: %.2f",

            totalPassengers, survivalRate, avgAge);

```



```

        result.set(resultString);
        context.write(key, result);
    }
}

// Main method to configure and run the MapReduce job
public static void main(String[] args) throws Exception {
    if (args.length < 2) {
        System.err.println("Usage: Aggregate <input file> <output dir>");
        System.exit(1);
    }

    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Titanic Aggregator");

    job.setJarByClass(Aggregate.class);
    job.setMapperClass(TitanicMapper.class);
    job.setReducerClass(TitanicReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(TitanicData.class);

    FileInputFormat.addInputPath(job, new Path(args[0])); // Input file path
    FileOutputFormat.setOutputPath(job, new Path(args[1])); // Output directory

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```